

KEYU HE

(213) 713-2973 | keyuhe@cmu.edu | [keyu-he.github.io](https://github.com/keyu-he) | www.linkedin.com/in/keyuhe/

Education

Carnegie Mellon University, Pittsburgh, PA

Master of Science in Intelligent Information Systems @ School of Computer Science

May 2027

GPA: 4.14/4.33

University of Southern California, Los Angeles, CA

Bachelor of Science in Computer Science & Bachelor of Art in Applied Mathematics

May 2025

GPA: 3.98/4.00

Minor in Artificial Intelligence Applications

Skills

Programming: C++, C, Python, Java, MySQL, HTML, CSS, JS, TS, x86-64 Assembly

Frameworks/Tools/Software: PyTorch, Pandas, NumPy, LangChain, FAISS, Git, AWS, L^AT_EX

Areas of Expertise: Machine Learning, Natural Language Processing (NLP), Large Language Models (LLMs), Vision Language Models (VLMs), Data Science / Data Engineering

Publications

1. **Keyu He**, Xuhui Zhou, Maarten Sap. *Social Gym and SPaRTan: Benchmarking and Improving LLM Social Reasoning via Multi-Agent Game Tournaments*. Preprint, 2026. [10.48550/arXiv.2608.09128](https://arxiv.org/abs/2608.09128)
[multi-agent environment engineering; LLM agent orchestration; large-scale benchmarking; self-play method design]
2. **Keyu He**, Tejas Srinivasan, Brihi Joshi, Xiang Ren, Jesse Thomason, Swabha Swayamdipta. *Believing without Seeing: Quality Scores for Contextualizing Vision-Language Model Explanations*. ACL 2026. [10.48550/arXiv.2509.25844](https://arxiv.org/abs/2509.25844)
[VLM metrics design; calibration; human-subject studies; statistical analyses]
3. **Keyu He***, Qianou Ma*, Valerie Chen, Wayne Chi, Tongshuang Wu. *RECAP: An End-to-End Platform for Capturing, Replaying, and Analyzing AI-Assisted Programming Interactions*. ACL Demo 2026. [10.48550/arXiv.2605.01104](https://arxiv.org/abs/2605.01104)
[VS Code extension; full-stack TypeScript development; cloud data pipeline design; analyses of human-AI interactions]
4. Brihi Joshi*, **Keyu He***, Sahana Ramnath, Sadra Sabouri, Kaitlyn Zhou, Souti Chattopadhyay, Swabha Swayamdipta, Xiang Ren. *ELI-Why: Evaluating the Pedagogical Utility of LLM Explanations*. Findings of ACL 2025. [10.48550/arXiv.2506.14200](https://arxiv.org/abs/2506.14200)
[Benchmark/Dataset design; human-subject studies; web-app implementation; multi-metric evaluation of LLM explanations]
5. Huihan Li*, Arnav Goel*, **Keyu He**, and Xiang Ren. *Attributing Culture-Conditioned Generations to Pretraining Corpora*. ICLR 2025. [10.48550/arXiv.2412.20760](https://arxiv.org/abs/2412.20760).
[large-scale corpus analysis; Python data engineering and visualization; survey server infrastructure design]

Experience

Adobe, San Jose, CA

May 2026 – Aug. 2026

Machine Learning Engineer Intern

- Shipped 5+ agent skills, including a four-skill family outside the team's own domain, through an **agent skill generator** merged into the team's shared skill repo, which turns rough developer requests into standards-compliant skills, grounding against internal prior art so existing work isn't rebuilt and self-auditing against the agentskills.io spec.
- Developed the systematic quality measure for the team's agent recommendation pipeline, a **reference-free 6-dimension rubric** combining LLM-as-judge with deterministic checks, reproducible to 0.01–0.05 per-dimension variance across repeated runs and surfacing two systematic weaknesses for the team to act on.
- Took that measure to production across the org: ported it onto the internal GenAI evaluation service as a reusable scorer and staged it end-to-end for two enterprise customers, on a four-stage Databricks pipeline built to run it.

CMind, Remote

May 2025 – Aug. 2025

Data Engineer Intern

- Built a LangChain-based RAG prototype over an internal analytics DB (chunking, embeddings, FAISS indexing/retrieval, LLM answer synthesis) with metadata filters and notebook/API utilities.
- Designed CEO-speech trait metrics and an auto-evaluation pipeline with custom rubrics, achieving Fleiss' κ of 0.89–0.91 against human annotations.